

MNIST Convnet Sample Code

Tutorial on MNIST CNN:

First Reference: <https://github.com/fchollet/deep-learning-with-python-notebooks/blob/master/2.1-a-first-look-at-a-neural-network.ipynb>

<https://github.com/fchollet/deep-learning-with-python-notebooks/blob/master/5.1-introduction-to-convnets.ipynb>

TF Keras has defined MNIST model and its training/testing image dataset, all you have to do is to import them into your program, which significantly save the time for learning. See github code below:

(1) The code for the handwritten digits recognition (to save trained model): 7-1convnets-NumeralDet-saveTrained.py

The folder: opencv/IP120-AI-DL/2018F/7-1convnets-NumeralDet-saveTrained.py /

The URL of the file: <https://github.com/hualili/opencv/blob/master/IP120-AI-DL/2018F/7-1convnets-NumeralDet-saveTrained.py>

Step 1 Understand the Architecture (MNIST Convnet)

model.summary

```
>>> model.summary()
```

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_1 (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_2 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_2 (MaxPooling2D)	(None, 5, 5, 64)	0
conv2d_3 (Conv2D)	(None, 3, 3, 64)	36928
flatten_1 (Flatten)	(None, 576)	0
dense_1 (Dense)	(None, 64)	36928
dense_2 (Dense)	(None, 10)	650
Total params: 93,322		
Trainable params: 93,322		
Non-trainable params: 0		

MNIST Convnet Architecture

model.summary

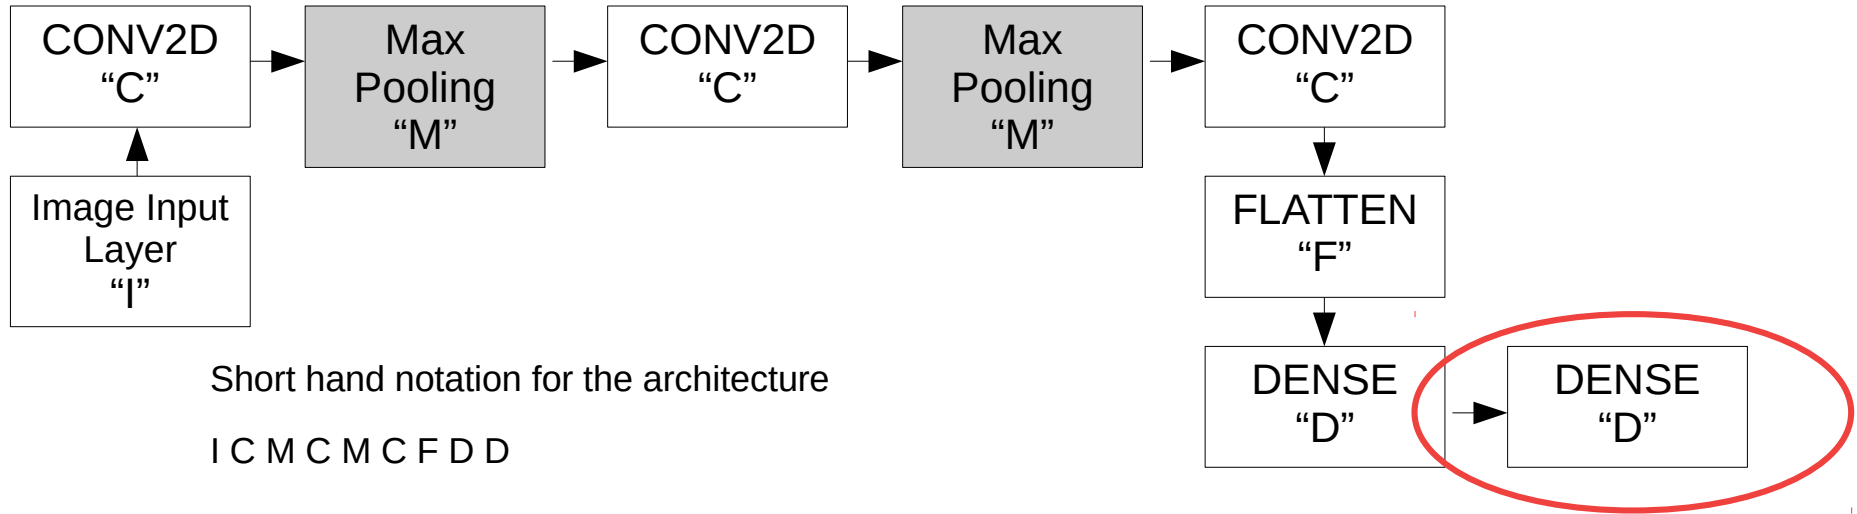
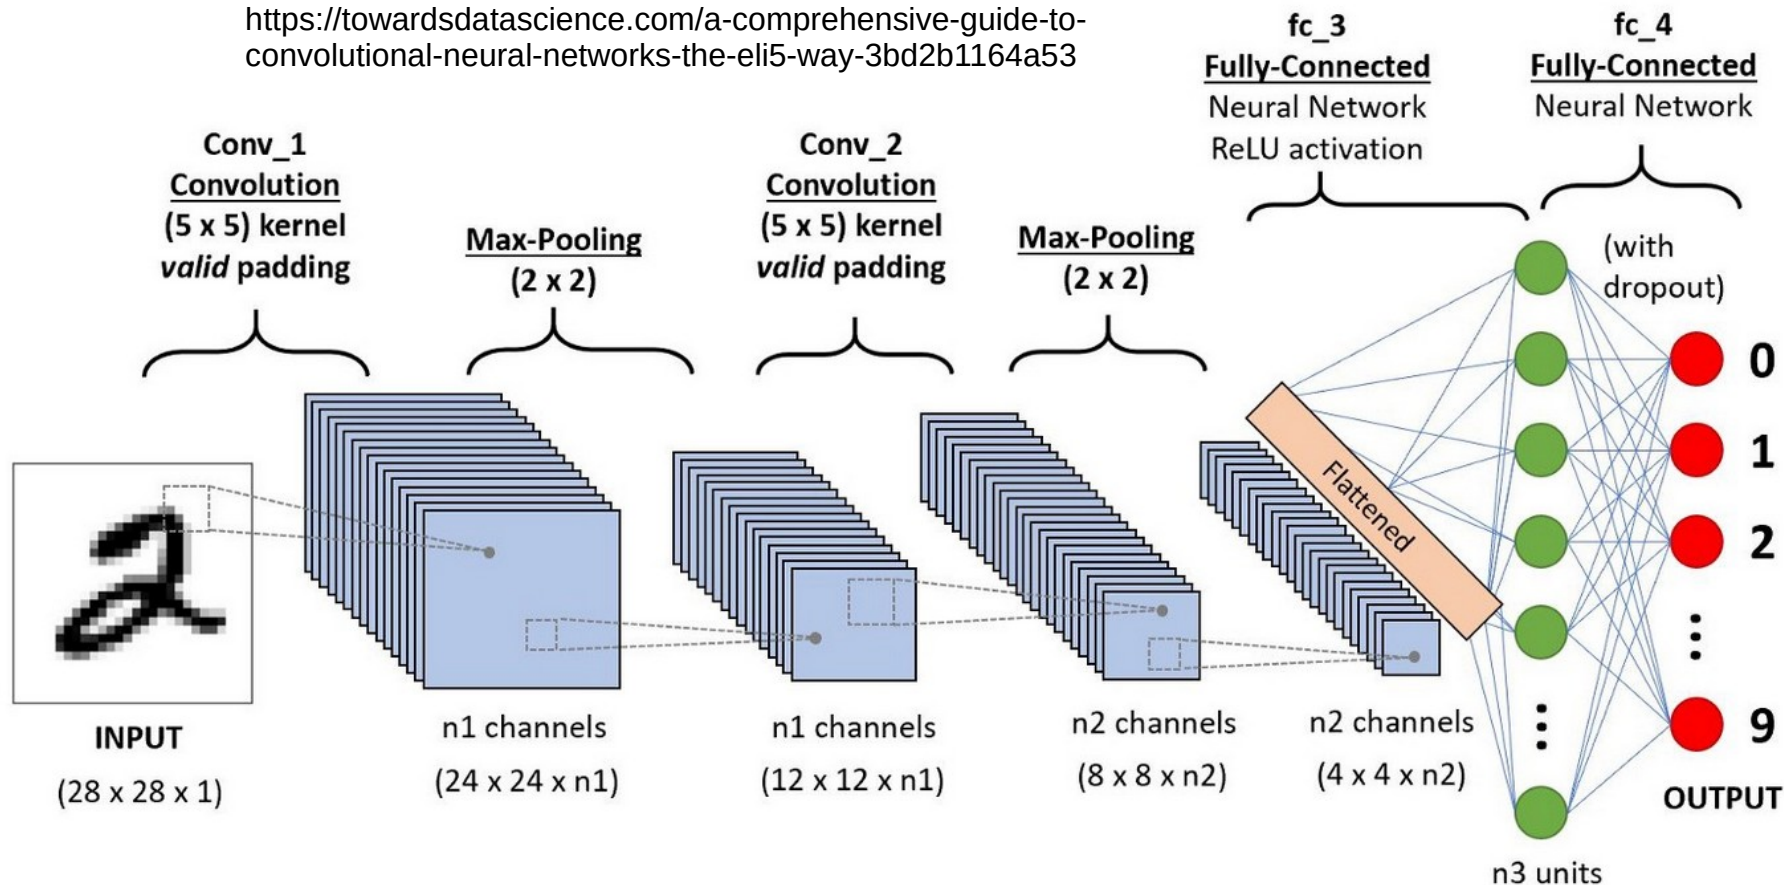



Illustration of A CNN for Digits Recognition

<https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53>



Convolution Neural Networks

<https://www.superdatascience.com/blogs/convolutional-neural-networks-cnn-summary>



2D Convolution

Reference for the theoretical background: Chapter 6, Robot Vision, pp. 104 – 111, by BKP Horn, MIT Press

Definition:


$$f(x) * g(x) = \int_{-\infty}^{\infty} f(\tau) \cdot g(x - \tau) d\tau$$

$$c(n_1, n_2) = \sum_{k_1=-\infty}^{\infty} \sum_{k_2=-\infty}^{\infty} a(k_1, k_2) b(n_1 - k_1, n_2 - k_2)$$

Image Kernel


Summation lower and upper bound in the case of M-by-N image $f(x,y)$, should be adjusted to $k_1 = 0$ to $M-1$, $k_2 = 0$ to $N-1$

Reference for the OpenCV implementation: Learning OpenCV, Chapter 6, pp. 144 – 164.



Note: (1) 3 primitive computations: shift, multiplication, and addition;
(2) use discrete 2D convolution formula to compute 5x5 sample image with 3x3 kernels

2D Convolution Example



[https://towardsdatascience.com/
simple-introduction-to-
convolutional-neural-networks-
cdf8d3077bac](https://towardsdatascience.com/simple-introduction-to-convolutional-neural-networks-cdf8d3077bac)

2D Convolution with Matlab/Octave

$C = \text{conv2}(A,B)$ computes the two-dimensional convolution of matrices A and B .

The size of C is determined as follows: if $[ma,na] = \text{size}(A)$, $[mb,nb] = \text{size}(B)$, Then $[mc,nc] = \text{size}(C)$, Where $mc = \max([ma+mb-1,ma,mb])$ and $nc = \max([na+nb-1,na,nb])$.



Octave
on Linux

```
>> A = [ 0 0 100 100 100  
        0 0 100 100 100  
        0 0 100 100 100  
        0 0 100 100 100  
        0 0 100 100 100 ]
```

```
>> B = [ 1 0 -1  
        1 0 -1  
        1 0 -1 ]
```

```
C = conv2(A,B)
```

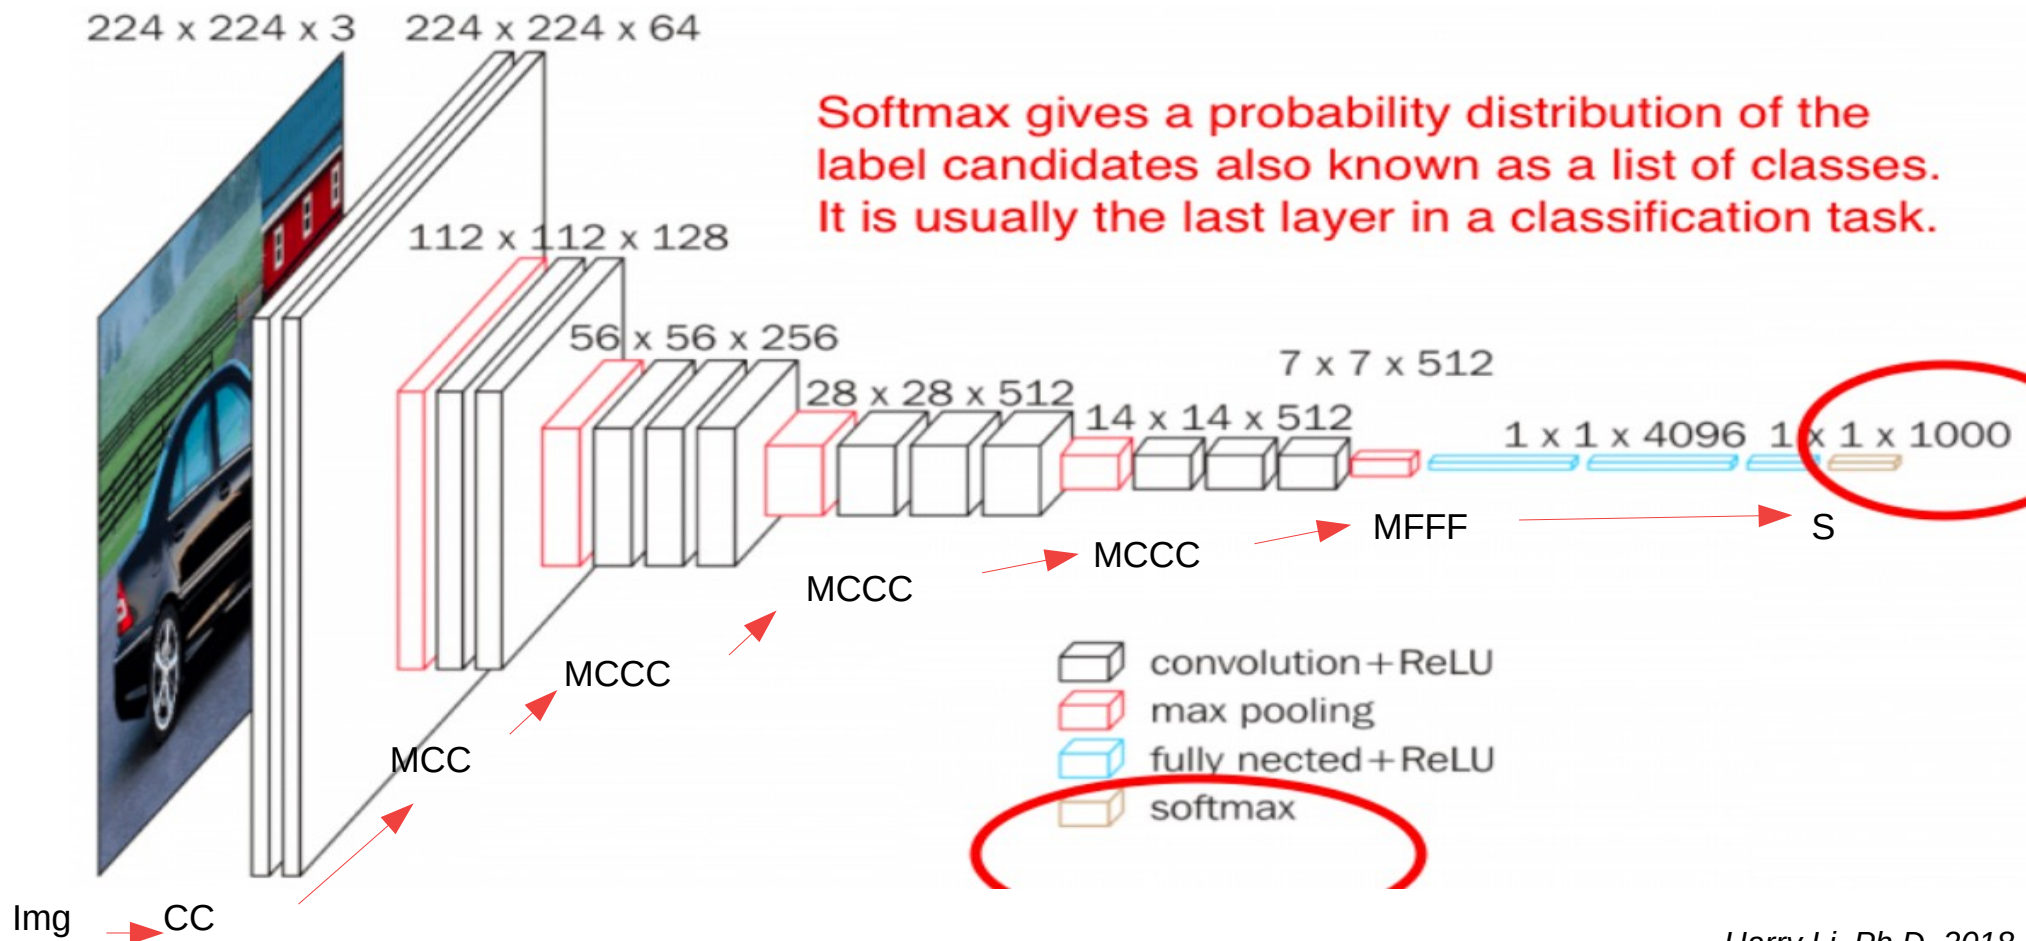
```
C =
```

```
0    0   100   100    0 -100 -100  
0    0   200   200    0 -200 -200  
0    0   300   300    0 -300 -300  
0    0   300   300    0 -300 -300  
0    0   300   300    0 -300 -300  
0    0   200   200    0 -200 -200  
0    0   100   100    0 -100 -100
```

<https://github.com/hualili/opencv/blob/master/deep-learning-2020S/1-2020S-%232019S-23-2DConvolution-2019-2-4.pdf>


Architecture Example VGG16

<https://medium.com/data-science-bootcamp/understand-the-softmax-function-in-minutes-f3a59641e86d>



Softmax Network

<https://medium.com/data-science-bootcamp/understand-the-softmax-function-in-minutes-f3a59641e86d>



The above Udacity lecture slide shows that Softmax function turns logits [2.0, 1.0, 0.1] into probabilities [0.7, 0.2, 0.1], and the probabilities sum to 1.

“logits layer is popularly used for the last neuron layer for classification task which produces raw prediction values as real numbers ranging from [-infinity, +infinity]”

“Softmax’s input is the output of the fully connected layer immediately preceding it, and it outputs the final output of the entire neural network. This output is a probability distribution of all the label class candidates.”

Output Layer


Example: NIST 10 digits (0-9) convnet Dense layer

```
from keras import models
from keras import layers
```


512 neurons in the
layer

```
network = models.Sequential()
network.add(layers.Dense(512, activation='relu', input_shape=(28 * 28,)))
network.add(layers.Dense(10, activation='softmax'))
```

The network consists of 2 Dense layers, densely-connected ("fully-connected") neural layers. The output layer is a 10-way "softmax" layer, it returns an array of 10 probability scores (summing to 1). Each score will be the probability that the current digit image belongs to one of our 10 digit classes.




Sample code :
<https://github.com/fchollet/deep-learning-with-python-notebooks>



*Illustration of the softmax block diagram from Milos Hauskrecht,
milos@cs.pitt.edu, 5329 Sennott Square*

Dense Layer



From: <https://towardsdatascience.com/exploring-activation-functions-for-neural-networks-73498da59b02>


Discussion with Feed Forward No Conv

Image Size:
28x28

768 input
(28x28)


512 neuron
Hidden layer

10 neuron output layer



From:
<https://towardsdatascience.com/exploring-activation-functions-for-neural-networks-73498da59b02>

3	9	12
9	9	7
22	22	17



Activation Functions

768 input
(28x28)

```
model = Sequential()  
model.add(Dense(512, input_shape=(784,)))  
model.add(Dense(10, activation='softmax'))
```

```
model = Sequential()  
model.add(Dense(512, activation='sigmoid', input_shape=(784,)))  
model.add(Dense(10, activation='softmax'))
```

```
model = Sequential()  
model.add(Dense(512, activation='tanh', input_shape=(784,)))  
model.add(Dense(10, activation='softmax'))
```

From: <https://towardsdatascience.com/exploring-activation-functions-for-neural-networks-73498da59b02>

Softmax

Example: NIST 10 digits (0-9) convnet Dense layer


```
from keras import models
from keras import layers
```

```
network = models.Sequential()
network.add(layers.Dense(512, activation='relu', input_shape=(28 * 28,)))
network.add(layers.Dense(10, activation='softmax'))
```

The network consists of 2 Dense layers, densely-connected ("fully-connected") neural layers. The output layer is a 10-way "softmax" layer, it returns an array of 10 probability scores (summing to 1). Each score will be the probability that the current digit image belongs to one of our 10 digit classes.



Sample code :
<https://github.com/fchollet/deep-learning-with-python-notebooks>



*Illustration of the softmax block diagram from Milos Hauskrecht,
milos@cs.pitt.edu, 5329 Sennott Square*


Prepare Image For Testing Trained MNIST convnet

Find out the input image size to MNIST convnet, from 2 sources: (1) from 7convnets-NumeralDetection-ch05.py (code below) (2) from the code, the dimension of the conv2d_1, kernel size is 3x3, from model.summary, the output shape is 26x26.

Example:

```
model = models.Sequential()  
model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)))
```

So, (1) use screen capture program to locate ROI as shown below, and then (2) we wrote opencv program to take any size input image and convert it to 28-by-



```
>>> model.summary()
```

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_1 (MaxPooling2)	(None, 13, 13, 32)	0
conv2d_2 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_2 (MaxPooling2)	(None, 5, 5, 64)	0
conv2d_3 (Conv2D)	(None, 3, 3, 64)	36928
flatten_1 (Flatten)	(None, 576)	0
dense_1 (Dense)	(None, 64)	36928
dense_2 (Dense)	(None, 10)	650
Total params: 93,322		
Trainable params: 93,322		
Non-trainable params: 0		

Convnet MNIST Architecture

Using TensorFlow backend.


Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_1 (MaxPooling2)	(None, 13, 13, 32)	0
conv2d_2 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_2 (MaxPooling2)	(None, 5, 5, 64)	0
conv2d_3 (Conv2D)	(None, 3, 3, 64)	36928

Total params: 55,744

Trainable params: 55,744

Non-trainable params: 0

10-29-2018 Road Map to Deploy CNN



10-29-2018 Save Keras Model

[https://www.google.com/search?](https://www.google.com/search?hl=en&ei=kFDZW__GM8SS0wLLhJSICg&q=save+model+json+keras&oq=save+Model+%28JSON%29&gs_l=psy-ab.1.0.0i22i30l7.4409.4409..7830...0.0..0.66.66.1.....0....1j2..gws-wiz.....0i71.fyoCDOFN8UU)

[hl=en&ei=kFDZW__GM8SS0wLLhJSICg&q=save+model+json+keras&oq=save+Model+%28JSON%29&gs_l=psy-ab.1.0.0i22i30l7.4409.4409..7830...0.0..0.66.66.1.....0....1j2..gws-wiz.....0i71.fyoCDOFN8UU](https://www.google.com/search?hl=en&ei=kFDZW__GM8SS0wLLhJSICg&q=save+model+json+keras&oq=save+Model+%28JSON%29&gs_l=psy-ab.1.0.0i22i30l7.4409.4409..7830...0.0..0.66.66.1.....0....1j2..gws-wiz.....0i71.fyoCDOFN8UU)


You can use `model.save(filepath)` to save a Keras model into a single HDF5 file which will contain: the architecture of the model, allowing to re-create the model. the weights of the model. the training configuration (loss, optimizer) the state of the optimizer, allowing to resume training exactly where you left off.

10-30-2018 Prepare h5py to Save Trained Network

<https://www.quora.com/How-do-I-save-a-convolution-neural-network-model-after-training-I-am-working-in-Python>

Keras for saving a model is just one line of code after training

Just make sure to have HDF5 for Python, use h5py package, which is a Pythonic interface to the HDF5 binary data format, which stores huge amounts of numerical data, and easily to be manipulated with NumPy. The files created can be exchanged including programs like IDL and MATLAB.



<http://www.h5py.org/>

First, install h5py, see the url link to the right and just follow the installation instructions, after the installation is done, be sure to use the following command to check:

```
import h5py
h5py.run_tests()
```

```
ubuntu@ubuntu-ThinkPad-Yoga-14:~/Documents$ python
Python 3.4.3 (default, Nov 28 2017, 16:41:13)
[GCC 4.8.4] on linux
Type "help", "copyright", "credits" or "license" for more information.
>>> import h5py
>>> h5py.run_tests()
.....X.....
.....X.....S..S...SS
.....SSSSSS
.....X...X.....X...X.....
-----
Ran 457 tests in 0.632s

OK (skipped=14, expected failures=6)
<unittest.runner.TextTestResult run=457 errors=0 failures=0>
>>>
```

10-30-2018 Save Trained Network with h5py


<https://www.quora.com/How-do-I-save-a-convolution-neural-network-model-after-training-I-am-working-in-Python>

Keras for saving a model is just one line of code after training

```
Model.save('harry_convnet_mnist.h5')
```

Example:

```
scores = model.evaluate(X_train,y_train,verbose=0)
print("Accuracy on train set: %.2f%%" % (scores[1]*100))
scores = model.evaluate(X_validation,y_validation,verbose=0)
print("Accuracy on validation set: %.2f%%" % (scores[1]*100))
scores = model.evaluate(X_test,y_test,verbose=0)
print("Accuracy on test set: %.2f%%" % (scores[1]*100))
model.save('ajits_model_d6_256.h5')
```



<http://www.h5py.org/>

Example: 7-1convnets-NumeralsDet-saveTrained.py

```
#-----save trained model-----*
import h5py
model.save('harryTest.h5')
#-end
```

Just add 2 lines of code

10-30-2018 Load Trained Network

<https://stackoverflow.com/questions/35074549/how-to-load-a-model-from-an-hdf5-file-in-keras>

Keras for loading a model is just one line of code

If you stored the complete model, not only the weights, in the HDF5 file, then load is as simple as

Example:

```
#-----load trained model-----*  
from keras.models import load_model  
model = load_model('harryTest.h5')  
model.summary() #check the model
```

```
>>> model.summary()
```

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_1 (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_2 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_2 (MaxPooling2D)	(None, 5, 5, 64)	0
conv2d_3 (Conv2D)	(None, 3, 3, 64)	36928
flatten_1 (Flatten)	(None, 576)	0
dense_1 (Dense)	(None, 64)	36928
dense_2 (Dense)	(None, 10)	650

```
Total params: 93,322  
Trainable params: 93,322  
Non-trainable params: 0
```

OpenCV Program To Resize 28x28 Images

7-3ResizeImage.py

```
import cv2
import numpy as np

gray = cv2.imread('test.png',cv2.IMREAD_GRAYSCALE)
gray = cv2.resize(gray, (28,28))
cv2.imshow('image',gray)
cv2.imwrite('resized_harryTest.jpg',gray, [int(cv2.IMWRITE_JPEG_QUALITY),90])

k = cv2.waitKey(0)
if k == 27:      # wait for ESC key to exit
    cv2.destroyAllWindows()
```

11-2-2018 Deploy MNIST with Test Image (1)

Example:

```
model = load_model('Numeral_detector.h5')
test_image= cv2.imread('resized_LCTest.jpg', cv2.IMREAD_GRAYSCALE)
test_image = np.reshape(test_image, [1, 28, 28, 1])
result_arr = model.predict(test_image)
result = model.predict_classes(test_image)
print(result_arr, ' ', result)
```


10-29-2018 Deploy MNIST with Test Image (1)

<https://medium.com/@o.kroeger/tensorflow-mnist-and-your-own-handwritten-digits-4d1cd32bbab4>

```
# create an array where to store 4 pics with one numeral each
```

```
# each image consists of total 784 pixels
```

```
images = np.zeros((4,784))
```

```
# and the correct values
```

```
correct_vals = np.zeros((4,10))
```

```
# test images 8, 0, 4, 3
```

```
i = 0
```

```
for no in [8,0,4,3]:
```

```
    gray = cv2.imread("img/blog/own_"+str(no)+".png", cv2.CV_LOAD_IMAGE_GRAYSCALE)
```

```
# resize the images and invert it (black background)
```

```
gray = cv2.resize(255-gray, (28, 28))
```

```
# save the processed images
```

```
cv2.imwrite("pro-img/image_"+str(no)+".png", gray)
```

```
"""
```


```
all images in the training set ranges from 0-1
```

```
not from 0-255 so divide the flatten images
```

```
(a one dimensional vector with 784 pixels)
```

```
"""
```

```
flatten = gray.flatten() / 255.0
```



8.png 0.png 4.png 3.png

Image size: 28x28 = 784

10-29-2018 Deploy MNIST with Test Image (2)

```
"""----- store flatten image and generate-----  
the correct_vals array  
correct_val for the first digit (9) would be  
[0,0,0,0,0,0,0,0,0,1]  
-----"""  
  
images[i] = flatten  
correct_val = np.zeros((10))  
correct_val[no] = 1  
correct_vals[i] = correct_val  
i += 1  
""" -----the prediction will be an array with four values-----  
which show the predicted number  
-----"""  
  
prediction = tf.argmax(y,1)  
  
"""-----run the prediction and the accuracy function-----  
using our generated arrays (images and correct_vals)  
-----"""  
print sess.run(prediction, feed_dict={x: images, y_: correct_vals})  
print sess.run(accuracy, feed_dict={x: images, y_: correct_vals})
```

11-2-2018 Architecture Aspects on Deployment

Homework:

1. USE Architecture of the Choice, and the Self written digits to Test the Pretrained Network.

2. Digitize 18 categories of Images from our Sample video (Show & tell in class Next week).

```
1
from keras import models
from keras import layers
from keras import optimizers

2
import matplotlib.pyplot as plt

plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()

plt.figure()

plt.show()
```

```
3
model = models.Sequential()
model.add(layers.Dense(256, activation='relu', input_dim=4 * 4 * 512))
model.add(layers.Dropout(0.5))
model.add(layers.Dense(1, activation='sigmoid'))
```

```
4
model.compile(optimizer=optimizers.RMSprop(lr=2e-5),
              loss='binary_crossentropy',
              metrics=['acc'])
```

11-2-2018 Architecture Aspects on Deployment

1

```
from keras import models  
from keras import layers
```

2

```
model = models.Sequential()  
model.add(conv_base)  
model.add(layers.Flatten())  
model.add(layers.Dense(256, activation='relu'))
```


11-6-2018 Keras Supported CNN Architectures

<https://www.pyimagesearch.com/2017/03/20/imagenet-vggnet-resnet-inception-xception-keras/>

Xception
InceptionV3
ResNet50
VGG16
VGG19
MobileNet

ResNet architecture which can be successfully trained at depths of 50-200 for ImageNet and over 1,000 for CIFAR-10

Two major drawbacks with VGGNet: (1) painfully slow to train. (2) Network architecture weights themselves are quite large (in terms of disk/bandwidth). Due to its depth and number of fully-connected nodes, VGG16 is over 533MB and VGG19 is 574MB. This makes deploying VGG a tiresome task.



Cornell University
Library

[arXiv.org](#) > [cs](#) > [arXiv:1603.05027](#)

[Computer Science](#) > [Computer Vision and Pattern Recognition](#)

Identity Mappings in Deep Residual Networks

[Kaiming He](#), [Xiangyu Zhang](#), [Shaoqing Ren](#), [Jian Sun](#)


(Submitted on 16 Mar 2016 (v1), last revised 25 Jul 2016 (this version, v3))

<https://github.com/KaimingHe/resnet-1k-layers>

11-6-2018 Use Keras Supported CNN

<https://www.pyimagesearch.com/2017/03/20/imagenet-vggnet-resnet-inception-xception-keras/>

```
# import the necessary packages
from keras.applications import ResNet50
from keras.applications import InceptionV3
from keras.applications import Xception # TensorFlow ONLY
from keras.applications import VGG16
from keras.applications import VGG19
from keras.applications import imagenet_utils
from keras.applications.inception_v3 import preprocess_input
from keras.preprocessing.image import img_to_array
from keras.preprocessing.image import load_img
import numpy as np
import argparse
import cv2
```



Note: Weights for VGG16 and VGG19 are > 500MB. ResNet weights are ~100MB, while Inception and Xception weights are between 90-100MB. If this is the first time you are running this script for a given network, these weights will be (automatically) downloaded and cached to your local disk. Depending on your internet speed, this may take awhile. However, once the weights are downloaded, they will not need to be downloaded again, allowing subsequent runs of `classify_image.py` to be much faster.


11-6-2018 Deploy Cats Dogs Detection CNN

```
import time
import keras
keras.__version__
import numpy as np

from keras import models
import cv2
from keras.models import load_model
model = load_model('cats_and_dogs_small_1.h5')
#model.summary() #check the model

image = cv2.imread('1.jpg', cv2.IMWRITE_JPEG_QUALITY)
image = cv2.resize(image, (150, 150))
image = np.reshape(image, [1, 150, 150, 3])

time_curr = int(round(time.time() * 1000))
output = model.predict_classes(image)
time_end = int(round(time.time() * 1000))
time = time_end - time_curr
print("")
print(output)
print("")
print(time)
print("")
```



One image with resolution
150x150 with 3 channels (r,
g, b) colors

11-9-2018 Python Imports for Augmentation To Create Train Data Set (1)

<https://leemendelowitz.github.io/blog/how-does-python-find-packages.html>

Python finds packages: Python imports by searching the directories listed in `sys.path`

```
>>> import sys
>>> print '\n'.join(sys.path)

/usr/lib/python2.7
/usr/lib/python2.7/plat-x86_64-linux-gnu
/usr/lib/python2.7/lib-tk
/usr/lib/python2.7/lib-old
/usr/lib/python2.7/lib-dynload
/usr/local/lib/python2.7/dist-packages
/usr/lib/python2.7/dist-packages
/usr/lib/python2.7/dist-packages/PILcompat
/usr/lib/python2.7/dist-packages/gtk-2.0
/usr/lib/python2.7/dist-packages/ubuntu-ssoclient
>>>
```

`sys.path.append()` adds the path for the code to access.

10-26-2018 Keras API Functions

```
1  
model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)))
```

```
2  
model.add(layers.MaxPooling2D((2, 2)))
```

```
3  
model.add(layers.Flatten())
```

```
4  
model.add(layers.Dense(64, activation='relu'))
```

```
5  
tf.keras.layers.Dropout(0.2)(fc_1)
```

```
6  
model.summary()
```

Oct-Nov-2018 Keras API Functions

1

```
import input_data
```

2

```
model.save('harryTest.h5')
```

3

```
from keras.models import load_model
```

4

```
model = load_model('harryTest.h5')
```

5

```
del model
```

6

```
import h5py  
h5py.run_tests()
```

7

```
result_arr = model.predict(test_image)
```