

Practical Python For Tensorflow

Tony Xu
10/7/2017



Agenda

- **Recap of what we have learned last time**
- Python Solution
- C/C++ Solution
- TensorFlow Solution



Practice on Convolution

- Assuming stride = 1; height and width are same for input, output, kernel.
- Q: if we want output size equals to the input size. What we need to do to the input?
- ->padding
- What's the size for padding, if kernel size = k?
- ->padding size = $(k-1)/2$, k=kernel size
- More general formula : $O = (I - k + 2p)/s + 1$
 - O is output size
 - I is input size
 - k is kernel size
 - p is padding size
 - s is stride size, usually is 1



Agenda

- Recap of what we have learned last time
- **Python Solution**
- C/C++ Solution
- TensorFlow Solution



Python Solution

- **We only use Python primitives to illustrate how to program in Python, but you could use Numpy (Numerical Python)...**
- **We are going to show you step by step with a few small examples before we form a solution.**



More on Python List

- Initialization of a short list
- `>>> a = [0, 0, 0]`
- How about a long list ...100 items?
1000?
- `>>> a = [0 for i in range(100)]`
- How about multiple dimensions?
- `>>> a = [[0 for i in range(10)] for j in range(5)]`

Rows

Cols



Demo



Agenda

- Recap of what we have learned last time
- Python Solution
- **C/C++ Solution**
- TensorFlow Solution



C/C++ Solution

- **Assume you know C/C++ already**
- **The challenge/difficulty is to access a 2D array with pointers**



Demo



Agenda

- Recap of what we have learned last time
- Python Solution
- C/C++ Solution
- **TensorFlow Solution**



Quiz

- **How many sections does a TensorFlow program usually have?**
- **2**
- **What are they?**
- **Graph and Session**
- **Why do we divide it to two sections?**
- **Reduce the overhead, easily distribute jobs to different devices.**



TensorFlow Solution

- **Introduce a few TensorFlow functions.**
- **Practices in Convolution**
- **We also show you a small step each time and finally reach to our goal.**



Tensor Shape and Reshape

- **We always need to manipulate tensor's shape to meet desired math**

```
# tensor 't' is [1, 2, 3, 4, 5, 6, 7, 8, 9]
# tensor 't' has shape [9]
```

```
reshape(t, [3, 3]) ==> [[1, 2, 3],
                          [4, 5, 6],
                          [7, 8, 9]]
```

```
# tensor 't' is [[[1, 1, 1],
#                  [2, 2, 2]],
#                [[3, 3, 3],
#                  [4, 4, 4]],
#                [[5, 5, 5],
#                  [6, 6, 6]]]
# tensor 't' has shape [3, 2, 3]
```

```
reshape(t, [-1]) ==> [1, 1, 1, 2, 2, 2, 3, 3, 3, 4, 4, 4, 5, 5, 5, 6, 6, 6]
```



TensorFlow: argmax

- **tf.argmax:** Returns the index with the largest value across axes of a tensor



```
pred = np.array([[31, 23, 4, 24, 27, 34],  
                 [18, 3, 25, 0, 6, 35],  
                 [28, 14, 33, 22, 20, 8],  
                 [13, 30, 21, 19, 7, 9],  
                 [16, 1, 26, 32, 2, 29],  
                 [17, 12, 5, 11, 10, 15]])
```

- **>>> argmax(pred, 1)**
- **array([5, 5, 2, 1, 3, 0])**
- **Remember: 0: vertical, 1: horizontal**



TensorFlow: reduce mean

- **tf.reduce_mean**: Computes the mean of elements across dimensions of a tensor.

```
# 'x' is [[1., 1.]  
#         [2., 2.]  
tf.reduce_mean(x) ==> 1.5  
tf.reduce_mean(x, 0) ==>   
tf.reduce_mean(x, 1) ==> 
```



TensorFlow: Equal

- **`tf.equal(x,y)`: Returns the truth value of $(x == y)$ element-wise. (Returns a true-false tensor)**
- **`x, y` are tensors with same shape.**
- **If we have `x = tf.constant([1,2,3])`, then `tf.equal(x,x)` will give us:**
- **`[true, true, true]`**
- **`tf.not_equal`, `tf.less` etc. are in similar fashion.**



Test Your Knowledge

```
# Test trained model  
correct_prediction = tf.equal(tf.argmax(y, 1), tf.argmax(y_, 1))  
accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))  
print(sess.run(accuracy, feed_dict={x: mnist.test.images,  
                                     y_: mnist.test.labels}))
```



TensorFlow: conv2d

- **It computes a 2-D convolution given 4-D input and filter tensors**
- **Why 4D?**

```
conv2d(  
    input,  
    filter,  
    strides,  
    padding,  
    use_cudnn_on_gpu=None,  
    data_format=None,  
    name=None  
)
```

A 4D Tensor with NHWC format:
[100,28,28,1]

A 4D Tensor with HWCC
format: [3,3,1,32]

A 1D Tensor with NHWC
format:[1,1,1,1]

"SAME", "VALID"

Default "NHWC"



Demo



Backup



Softmax

In [mathematics](#), the softmax function, or normalized exponential function,^{[1]:198} is a generalization of the [logistic function](#) that "squashes" a K -dimensional vector \mathbf{z} of arbitrary real values to a K -dimensional vector $\sigma(\mathbf{z})$ of real values in the range $[0, 1]$ that add up to 1. The function is given by

$$\sigma(\mathbf{z})_j = \frac{e^{z_j}}{\sum_{k=1}^K e^{z_k}} \quad \text{for } j = 1, \dots, K. \quad \text{"Euler's number" (2.71828)}$$

$$\mathbf{z} = [1.0, 2.0, 3.0, 4.0, 1.0, 2.0, 3.0]$$

$$= e^{z_j} [2.72, 7.39, 20.09, 54.6, 2.72, 7.39, 20.09]$$

$$\sum_{k=1}^K e^{z_k} = 114.8$$

Add up to 1

$$\sigma(\mathbf{z})_j = [0.024, 0.064, 0.175, 0.475, 0.024, 0.064, 0.175]$$



How to use Reshape

- `input_layer = tf.reshape(features["x"], [-1, 28, 28, 1])`
- Note that we've indicated -1 for batch size, which specifies that this dimension should be dynamically computed based on the number of input values in `features["x"]`, holding the size of all other dimensions constant. This allows us to treat `batch_size` as a hyperparameter that we can tune. For example, if we feed examples into our model in batches of 5, `features["x"]` will contain 3,920 values (one value for each pixel in each image), and `input_layer` will have a shape of `[5, 28, 28, 1]`. Similarly, if we feed examples in batches of 100, `features["x"]` will contain 78,400 values, and `input_layer` will have a shape of `[100, 28, 28, 1]`.



Cost Function

- `tf.nn.softmax_cross_entropy_with_logits`
- Format:

```
softmax_cross_entropy_with_logits(  
    _sentinel=None,  
    labels=None,  
    logits=None,  
    dim=-1,  
    name=None  
)
```

- Computes softmax cross entropy between logits and labels



The End

