

11-13-2018 SqueezeDet for Pedestrian Detection

<https://medium.com/omnius/fast-object-detection-with-squeezedet-on-keras-5cdd124b46ce>

Single-shot Object Detection:

(1) identify the locations, usually boxes, and the classes of multiple objects within the image.

(2) provide with a confidence value of its prediction for filtering.

Define anchor boxes: before the training, define a grid over the image. At each point of this grid, define multiple boxes with different aspect ratios tailored to the dataset.

<https://www.amazon.com/EVGA-GeForce-GAMING-Graphics-12G-P4-2990-KR/dp/B00UVN21RQ>



TITAN X GPU

EVGA GeForce GTX TITAN X 12GB
GAMING, Play 4k with Ease Graphics
Card 12G-P4-2990-KR

Buy used:
\$673.08
+ \$8.42 shipping

Harry Li, Ph.D. 2018

SqueezeDet: Unified, Small, Low Power Fully Convolutional Neural Networks for Real-Time Object Detection for Autonomous Driving

Bichen Wu¹, Alvin Wan¹, Forrest Iandola^{1,2}, Peter H. Jin¹, Kurt Keutzer^{1,2}

¹UC Berkeley, ²DeepScale

{bichen, alvinwan, phj, keutzer}@berkeley.edu, forrest@deepscale.ai

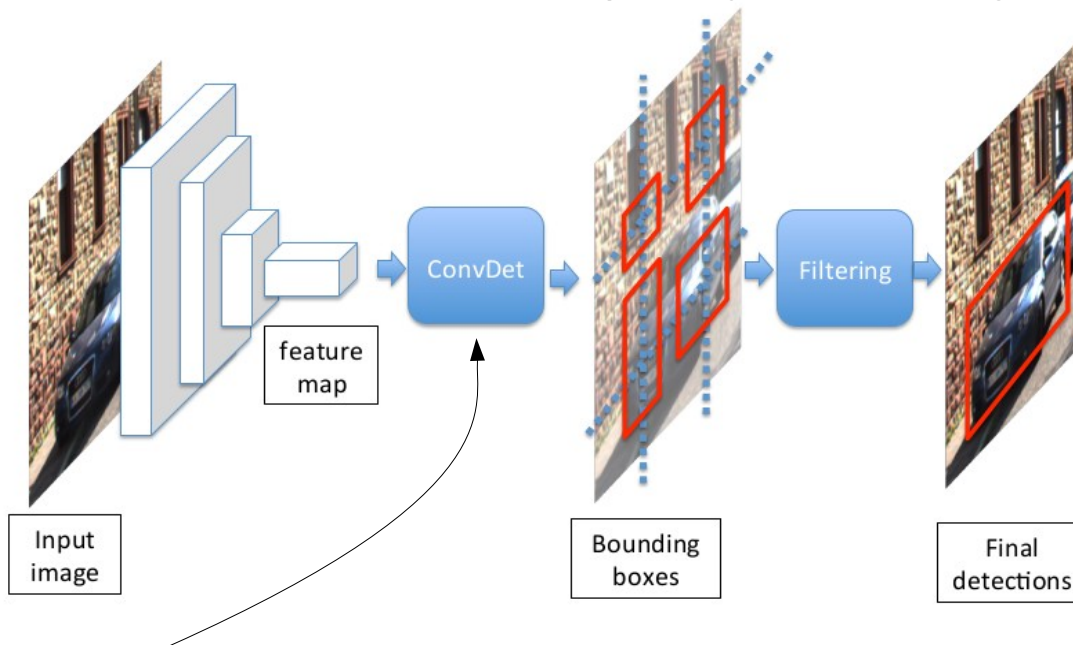
(1) Total model size is less than 8MB.

(2) The inference speed: 57.2 FPS 3 with input image resolution of 1242x375.

by [21]: first, we use stacked convolution filters to extract a high dimensional, low resolution feature map for the input image. Then, we use *ConvDet*, a convolutional layer to take the feature map as input and compute a large amount of object bounding boxes and predict their categories. Finally, we filter these bounding boxes to obtain final detections. The “backbone” convolutional neural net (CNN) architecture of our network is SqueezeNet [16], which achieves AlexNet level imageNet accuracy with a model size of $< 5\text{MB}$ that

11-13-2018 SqueezeDet Detection Pipeline

Inspired by YOLO, a single-stage detection pipeline



SqueezeDet: Unified, Small, Low Power Fully Convolutional Neural Networks for Real-Time Object Detection for Autonomous Driving
Bichen Wu ¹ , Alvin Wan ¹ , Forrest Iandola ^{1,2} , Peter H. Jin ¹ , Kurt Keutzer ^{1,2}
¹ UC Berkeley, ² DeepScale

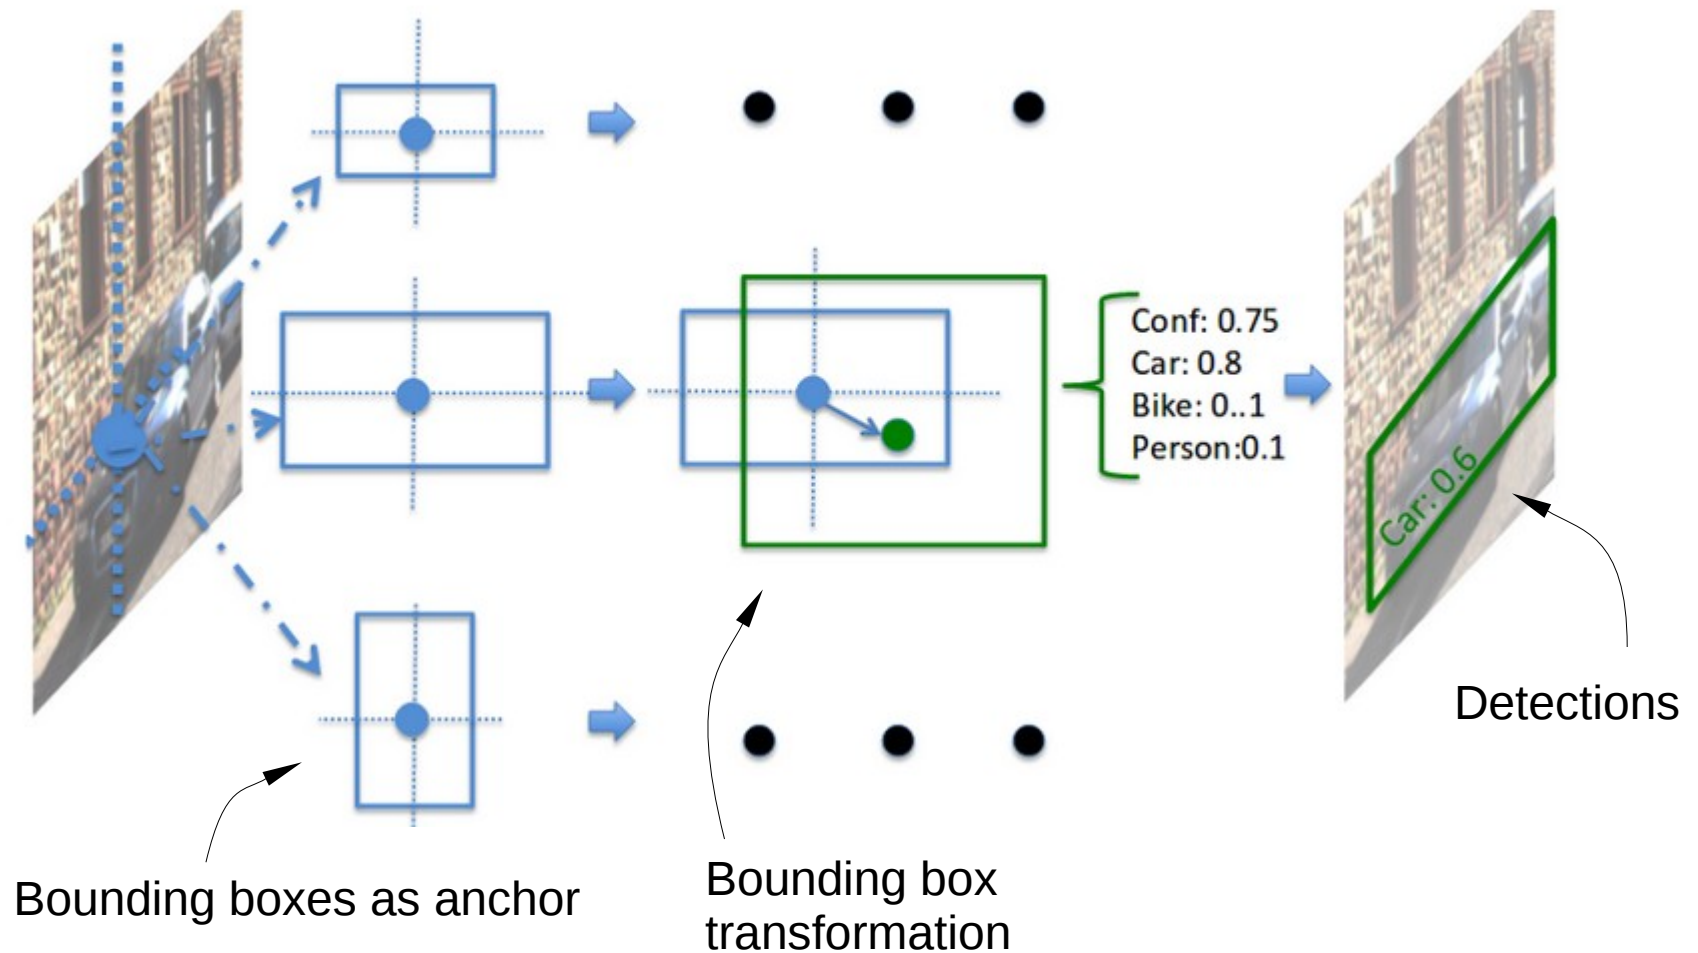
ConvDet is essentially a convolutional layer that is trained to output bounding box coordinates and class probabilities. It works as a sliding window that moves through each spatial position on the feature map.

$$\max_c Pr(\text{class}_c | \text{Object}) * Pr(\text{Object}) * \text{IOU}_{\text{truth}}^{\text{pred}} \quad (1)$$

Confidence of bounding box prediction


bounding boxes as anchor

11-13-2018 Bounding Box Transformation



SqueezeDet: Unified, Small, Low Power Fully Convolutional Neural Networks for Real-Time Object Detection for Autonomous Driving
Bichen Wu¹, Alvin Wan¹, Forrest Iandola^{1,2}, Peter H. Jin¹, Kurt Keutzer^{1,2}
¹ UC Berkeley, ² DeepScale


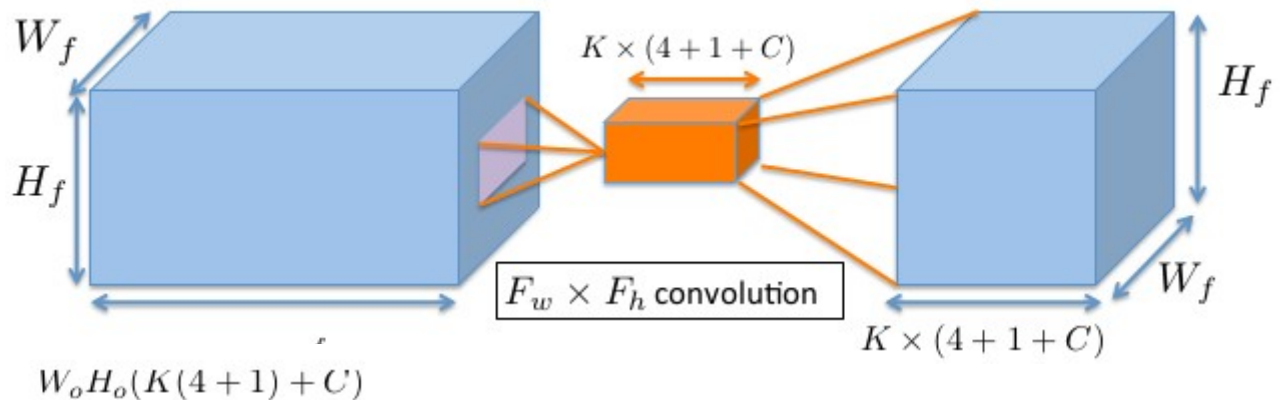
11-13-2018 Squeeze Net vs YOLO Architecture



SqueezeDet: Unified, Small, Low Power Fully Convolutional Neural Networks for Real-Time Object Detection for Autonomous Driving
 Bichen Wu¹, Alvin Wan¹, Forrest Iandola^{1,2}, Peter H. Jin¹, Kurt Keutzer^{1,2}
¹ UC Berkeley, ² DeepScale

(a) Last layer of Region Proposal Network (RPN) is a 1×1 convolution with $K \times (4 + 1)$ outputs.

(b) The ConvDet layer is a $F_w \times F_h$ convolution with output size of $K \times (5 + C)$.



(c) The detection layer of YOLO [21] contains 2 fully connected layers. The first one is of size $W_f H_f C$. The second one is of size $F_{fc1} W_o H_o K(5 + C)$.

11-13-2018 Performance Comparison and github Repo

method	E	car M	H	E	cyclist M	H	E	pedestrian M	H	mAP	model size (MB)	speed (FPS)
SubCNN [26]	90.8	89.0	79.3	79.5	71.1	62.7	83.3	71.3	66.4	77.0	-	0.2
MS-CNN [3]	90.0	89.0	76.1	84.1	75.5	66.1	83.9	73.7	68.3	78.5	-	2.5
PNET*	81.8	83.6	74.2	74.3	58.6	51.7	77.2	64.7	60.4	69.6	-	10
Pie*	89.4	89.2	74.2	84.6	76.3	67.6	84.9	73.2	67.6	78.6	-	0.83
FRCN+VGG16 [2]	92.9	87.9	77.3	-	-	-	-	-	-	-	485	1.7
FRCN+Alex [2]	94.7	84.8	68.3	-	-	-	-	-	-	-	240	2.9
SqueezeDet (ours)	90.2	84.7	73.9	82.9	75.4	72.1	77.1	68.3	65.8	76.7	7.9	57.2
SqueezeDet+ (ours)	90.4	87.1	78.9	87.6	80.3	78.1	81.4	71.3	68.5	80.4	26.8	32.1
VGG16 + <i>ConvDet</i> (ours)	93.5	88.1	79.2	85.2	78.4	75.2	77.9	69.1	65.1	79.1	57.4	16.6
ResNet50 + <i>ConvDet</i> (ours)	92.9	87.9	79.4	85.0	78.5	76.6	67.3	61.6	55.6	76.1	35.1	22.5

<https://github.com/omni-us/squeezedet-keras>

Reference: SqueezeDet: Unified, Small, Low Power Fully Convolutional Neural Networks for Real-Time Object Detection for Autonomous Driving
Bichen Wu¹, Alvin Wan¹, Forrest Iandola^{1,2}, Peter H. Jin¹, Kurt Keutzer^{1,2}
¹ UC Berkeley, ² DeepScale

Github code installation guide

<https://github.com/omni-us/squeezedet-keras/blob/master/Install.md>

10-30-2018 Save Trained Network with h5py


<https://www.quora.com/How-do-I-save-a-convolution-neural-network-model-after-training-I-am-working-in-Python>

Keras for saving a model is just one line of code after training

```
Model.save('harry_convnet_mnist.h5')
```

Example:

```
scores = model.evaluate(X_train,y_train,verbose=0)
print("Accuracy on train set: %.2f%%" % (scores[1]*100))
scores = model.evaluate(X_validation,y_validation,verbose=0)
print("Accuracy on validation set: %.2f%%" % (scores[1]*100))
scores = model.evaluate(X_test,y_test,verbose=0)
print("Accuracy on test set: %.2f%%" % (scores[1]*100))
model.save('ajits_model_d6_256.h5')
```



<http://www.h5py.org/>

Example: 7-1convnets-NumeralDet-saveTrained.py

```
#-----save trained model-----*
import h5py
model.save('harryTest.h5')
#-end
```

Just add 2
lines of code

10-30-2018 Load Trained Network

<https://stackoverflow.com/questions/35074549/how-to-load-a-model-from-an-hdf5-file-in-keras>

Keras for loading a model is just one line of code

If you stored the complete model, not only the weights, in the HDF5 file, then load is as simple as

Example:

```
#-----load trained model-----*  
from keras.models import load_model  
model = load_model('harryTest.h5')  
model.summary() #check the model
```

```
>>> model.summary()
```

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_1 (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_2 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_2 (MaxPooling2D)	(None, 5, 5, 64)	0
conv2d_3 (Conv2D)	(None, 3, 3, 64)	36928
flatten_1 (Flatten)	(None, 576)	0
dense_1 (Dense)	(None, 64)	36928
dense_2 (Dense)	(None, 10)	650
Total params: 93,322		
Trainable params: 93,322		
Non-trainable params: 0		

11-6-2018 Keras Supported CNN Architectures

<https://www.pyimagesearch.com/2017/03/20/imagenet-vggnet-resnet-inception-xception-keras/>

Xception
InceptionV3
ResNet50
VGG16
VGG19
MobileNet

ResNet architecture which can be successfully trained at depths of 50-200 for ImageNet and over 1,000 for CIFAR-10

Two major drawbacks with VGGNet: (1) painfully slow to train. (2) Network architecture weights themselves are quite large (in terms of disk/bandwidth). Due to its depth and number of fully-connected nodes, VGG16 is over 533MB and VGG19 is 574MB. This makes deploying VGG a tiresome task.



Cornell University
Library

[arXiv.org](#) > [cs](#) > [arXiv:1603.05027](#)

[Computer Science](#) > [Computer Vision and Pattern Recognition](#)

Identity Mappings in Deep Residual Networks

[Kaiming He](#), [Xiangyu Zhang](#), [Shaoqing Ren](#), [Jian Sun](#)


(Submitted on 16 Mar 2016 (v1), last revised 25 Jul 2016 (this version, v3))

<https://github.com/KaimingHe/resnet-1k-layers>

11-6-2018 Use Keras Supported CNN

<https://www.pyimagesearch.com/2017/03/20/imagenet-vggnet-resnet-inception-xception-keras/>

```
# import the necessary packages
from keras.applications import ResNet50
from keras.applications import InceptionV3
from keras.applications import Xception # TensorFlow ONLY
from keras.applications import VGG16
from keras.applications import VGG19
from keras.applications import imagenet_utils
from keras.applications.inception_v3 import preprocess_input
from keras.preprocessing.image import img_to_array
from keras.preprocessing.image import load_img
import numpy as np
import argparse
import cv2
```



Note: Weights for VGG16 and VGG19 are > 500MB. ResNet weights are ~100MB, while Inception and Xception weights are between 90-100MB. If this is the first time you are running this script for a given network, these weights will be (automatically) downloaded and cached to your local disk. Depending on your internet speed, this may take awhile. However, once the weights are downloaded, they will not need to be downloaded again, allowing subsequent runs of `classify_image.py` to be much faster.

11-6-2018 Deploy Cats Dogs Detection CNN

```
import time
import keras
keras.__version__
import numpy as np

from keras import models
import cv2
from keras.models import load_model
model = load_model('cats_and_dogs_small_1.h5')
#model.summary() #check the model

image = cv2.imread('1.jpg',cv2.IMWRITE_JPEG_QUALITY)
image = cv2.resize(image,(150,150))
image=np.reshape(image,[1,150,150,3])

time_curr=int(round(time.time()*1000))
output=model.predict_classes(image)
time_end=int(round(time.time()*1000))
time=time_end-time_curr
print("")
print(output)
print("")
print(time)
print("")
```

One image with resolution 150x150 with 3 channels (r, g, b) colors

11-10-2018 Amazon Rekognition


Amazon Rekognition

Easily add intelligent image and video analysis to your applications.

<https://aws.amazon.com/rekognition/>

Submit an image or video to the Rekognition API, and the service can identify the objects, people, text, scenes, and activities, as well as detect any inappropriate content, compare faces for a wide variety of user verification, people counting, and public safety use cases

1



Identify a person in a photo or video using your private repository of face images.

2



Pathing: capture the path of people in the scene, use the movement of athletes during a game to identify plays for post-game analysis.

3

Unsafe content detection, identify potentially unsafe or inappropriate content